Sources and further reading

Cited dimensions and identifiers were checked against papers and model cards. Illustrative values are invented.

Architecture

Vaswani et al. 2017. Attention Is All You Need. arXiv:1706.03762
He et al. 2016. Deep Residual Learning for Image Recognition. arXiv:1512.03385

Tokens, embeddings, output

Sennrich, Haddow, Birch 2016. Neural Machine Translation of Rare Words with Subword Units. arXiv:1508.07909
Mikolov et al. 2013. Efficient Estimation of Word Representations in Vector Space. arXiv:1301.3781
Press, Wolf 2017. Using the Output Embedding to Improve Language Models. arXiv:1608.05859

Position, normalization, activation

Su et al. 2021. RoFormer: Enhanced Transformer with Rotary Position Embedding. arXiv:2104.09864
Ba, Kiros, Hinton 2016. Layer Normalization. arXiv:1607.06450
Zhang, Sennrich 2019. Root Mean Square Layer Normalization. arXiv:1910.07467
Xiong et al. 2020. On Layer Normalization in the Transformer Architecture. arXiv:2002.04745
Hendrycks, Gimpel 2016. Gaussian Error Linear Units (GELUs). arXiv:1606.08415
Shazeer 2020. GLU Variants Improve Transformer. arXiv:2002.05202

What the parts learn

Elhage et al. 2021. A Mathematical Framework for Transformer Circuits. transformer-circuits.pub
Clark et al. 2019. What Does BERT Look At? arXiv:1906.04341
Tenney, Das, Pavlick 2019. BERT Rediscovers the Classical NLP Pipeline. arXiv:1905.05950
Geva et al. 2021. Transformer Feed-Forward Layers Are Key-Value Memories. arXiv:2012.14913
Olsson et al. 2022. In-context Learning and Induction Heads. arXiv:2209.11895

Decoding and efficiency

Holtzman et al. 2020. The Curious Case of Neural Text Degeneration. arXiv:1904.09751
Dao et al. 2022. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. arXiv:2205.14135
Ainslie et al. 2023. GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints. arXiv:2305.13245

Models cited for dimensions

Radford et al. 2019. Language Models are Unsupervised Multitask Learners (GPT-2). Config and code: github.com/openai/gpt-2
Brown et al. 2020. Language Models are Few-Shot Learners (GPT-3). arXiv:2005.14165
Touvron et al. 2023. Llama 2: Open Foundation and Fine-Tuned Chat Models. arXiv:2307.09288
Dubey et al. 2024. The Llama 3 Herd of Models. arXiv:2407.21783