Q01 Discussion Questions

1. What evidence would distinguish memorizing an answer from generalizing a rule to a new instance?
2. When does a chain-of-thought-like continuation count as evidence of a useful computation, and when might it be post-hoc text?
3. Which regularities in code are useful for next-token prediction but insufficient for software correctness?
4. How could a benchmark metric create the appearance of a sudden capability from a smooth underlying improvement?
5. If a model follows an instruction reliably, what additional evidence would you require before calling it an agent?