Q04 Discussion Questions

1. When is a final-answer verifier sufficient, and when is process supervision needed?
2. What reward loophole could a coding benchmark accidentally create?
3. How should an evaluation control for extra inference-time compute?
4. Can a polished rationale count as evidence of the computation that selected an answer? Why or why not?
5. What behavior would you refuse to optimize with an easily gamed proxy reward?