Q08 Discussion Questions

1. Which metric should a lab forecast before committing a large training budget?
2. What regime change would make a prior scaling fit least trustworthy?
3. How would you distinguish a benchmark ceiling from a true capability plateau?
4. Which meaning of forgetting is most relevant to a continually updated deployed model?
5. When is it rational to spend compute on data quality rather than parameters?