Q08 Deeper Project Ideas

1. Fit a scaling curve to matched small model runs, hold out one scale, and report prediction intervals.
2. Compare compute allocations across model size and token count on a small corpus.
3. Build a continual-learning test with before/after task scores and analyze forgetting by example type.