Q16 Discussion Questions

1. Which disclosed facts would you need before calling a closed model “non-transformer”?
2. Why can total parameters mislead when comparing MoE with dense models?
3. Which systems change could improve latency without changing learned behavior?
4. What long-context test would distinguish a large window from reliable use of distant evidence?
5. What matched experiment would make an alternative-architecture comparison credible?