Q17 Discussion Questions

1. Why is “transformer” insufficient as a performance explanation?
2. Which measurement distinguishes total MoE parameters from per-token cost?
3. How would you test whether a long context is actually used reliably?
4. When might a state-space model be a sensible candidate rather than attention?
5. Which controls are required for a fair architecture comparison?