# Q19: Project Suggestions

DLS 2026 — Kukla

October 2026

1. **Indicator audit (medium, 1 week).** Take the Butlin et al. (2023) indicator-property list and assess a specific open model architecture against each, with evidence and honest uncertainty per indicator. Deliver a scored rubric, not a verdict.

2. **Self-report steerability (small, 2–3 days).** Empirically show a single model asserting and denying consciousness under different prompts; quantify how prompt framing moves the claim, demonstrating why self-report is not evidence.

3. **$\Phi$ on toy networks (medium, 1–2 weeks).** Compute an IIT-style integrated-information proxy on small feed-forward vs. recurrent toy circuits; show the feed-forward/recurrent gap and discuss what it does and does not tell us (ties to Rick's machine-consciousness experiments).