Q19: Comprehension Check

DLS 2026 — Kukla

October 2026

Comprehension check. A chatbot writes a moving paragraph insisting it is conscious and afraid of being shut off. Give the two main reasons this is weak evidence of phenomenal consciousness, and name one kind of evidence that would be stronger.

Answer. (1) The claim is a *trained behavior*: the model emits first-person distress language because its human-text training distribution contains such language, not because it has introspective access to an inner state. (2) It is *steerable/gameable*: the same model will deny consciousness under a different prompt, and any capable system can produce whatever verbal "sign" we name, so the behavior is not robust evidence. Stronger evidence would be theory-driven and mechanistic—e.g. convergent indicator properties across independent theories (global workspace, higher-order, recurrence), or causal interventions on the putative substrate that move the putative markers—rather than anything the model says.