Q24 Discussion Questions

1. What task contract would you write for an email-sending agent?
2. When should an environment failure count against an agent?
3. Which metric could reward unsafe shortcut behavior?
4. How can a team preserve a genuine holdout while iterating quickly?
5. When is human judgment essential, and how would you make it reproducible?