Q28 Discussion

1. What benign explanation could cause a context-dependent refusal?
2. Why is an elicited reasoning trace not a direct mechanism readout?
3. Which control best distinguishes a cue from changed task content?
4. What safety measure is useful even without evidence of deception?
5. How should a team communicate uncertainty after an anomalous result?