## T1-01: Five discussion questions

1. When does a benchmark stop being evidence for a deployment decision?
2. Which AI-system handoff most often lacks an accountable owner?
3. Should a retrieval failure be attributed to the model? Defend a boundary.
4. What monitoring metric could improve while user harm rises?
5. Which decisions should remain human-authorized even with high measured accuracy?