T1-02: Five discussion questions

1. When is a learned representation preferable to a hand-engineered feature?
2. What task values are hidden inside a chosen loss?
3. Can an accurate gradient lead to a harmful model? Give a mechanism.
4. How would you detect leakage that survives a random train/test split?
5. Which error categories should override an improved aggregate metric?